

# Instant Neural Graphics Primitives with a Multiresolution Hash Encoding

THOMAS, NVIDIA, Switzerland  
ALEX EVANS, United Kingdom  
CHRISTOPH, NVIDIA, USA  
ALEXANDER, NVIDIA, Germany

Presented By Heguang Lin at 09/13/2023

# Motivation

Need a fast, compact representations with high-frequency detail for Multi-layer perceptrons (MLPs)

## Existing Methods


- Most efficient encodings are trainable, task-specific structures.
- Rely on heuristics, may complicate training, and may limit performance on GPUs.

## Proposed method:

- Adaptive, efficient, and task-independent.
- Achieves top-tier quality across tasks after brief training.
- Adaptivity: No structural updates needed during training.
- Efficiency: Hash table lookups are  $O(1)$ , and allows parallel querying for all resolutions.

# Existing Methods

(a) No encoding



411 k + 0 parameters  
11:28 (mm:ss) / PSNR 18.56

(b) Frequency  
[Mildenhall et al. 2020]



438 k + 0  
12:45 / PSNR 22.90

(c) Dense grid  
Single resolution



10 k + 33.6 M  
1:09 / PSNR 22.35

(d) Dense grid  
Multi resolution



10 k + 16.3 M  
1:26 / PSNR 23.62

(e) Hash table (ours)  
 $T = 2^{14}$



10 k + 494 k  
1:48 / PSNR 22.61

(f) Hash table (ours)  
 $T = 2^{19}$




10 k + 12.6 M  
1:47 / PSNR 24.58

Without encoding, the reconstruction is blurry.

# Existing Methods: Frequency Encoding

(a) No encoding



411 k + 0 parameters  
11:28 (mm:ss) / PSNR 18.56

(b) Frequency  
[Mildenhall et al. 2020]



438 k + 0  
12:45 / PSNR 22.90

(c) Dense grid  
Single resolution



10 k + 33.6 M  
1:09 / PSNR 22.35

(d) Dense grid  
Multi resolution



10 k + 16.3 M  
1:26 / PSNR 23.62

(e) Hash table (ours)  
 $T = 2^{14}$



10 k + 494 k  
1:48 / PSNR 22.61

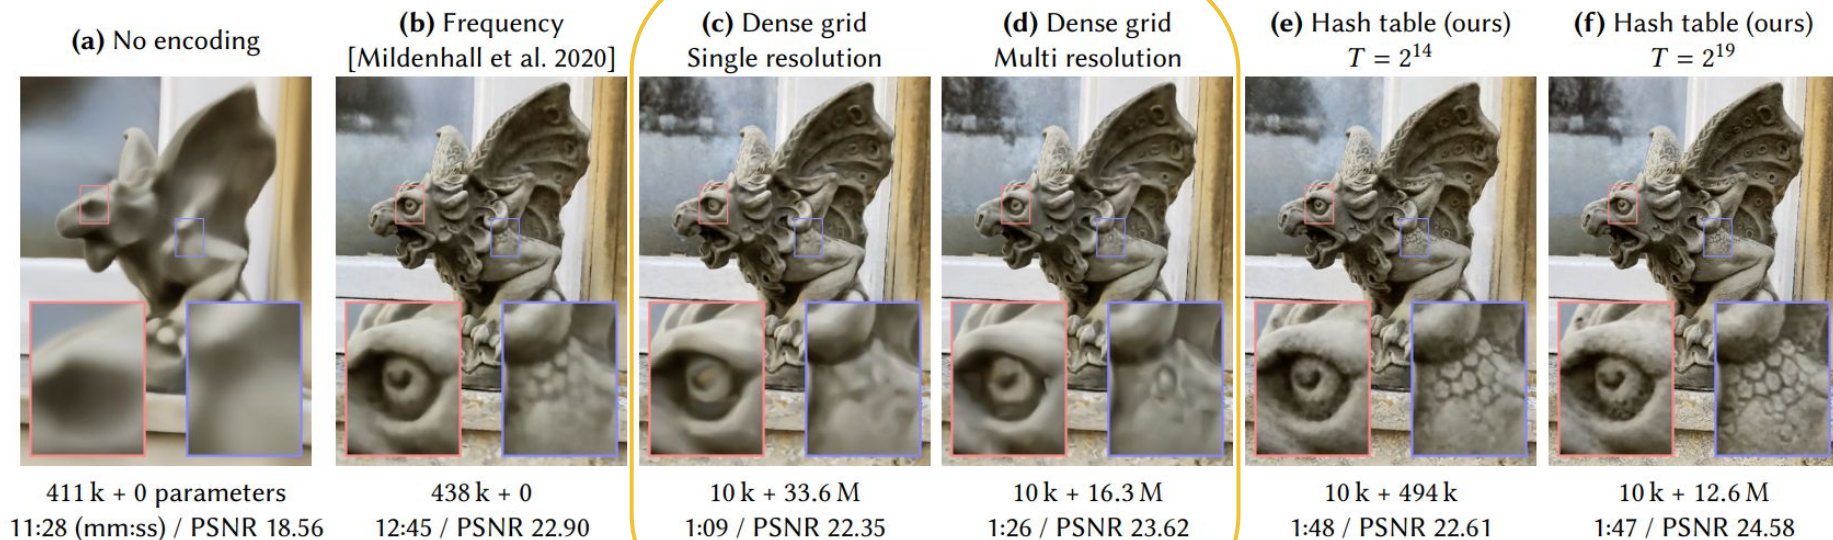
(f) Hash table (ours)  
 $T = 2^{19}$



10 k + 12.6 M  
1:47 / PSNR 24.58

$$\text{enc}(x) = ( \sin(2^0 x), \sin(2^1 x), \dots, \sin(2^{L-1} x), \\ \cos(2^0 x), \cos(2^1 x), \dots, \cos(2^{L-1} x) ) .$$

# Existing Methods: Parametric Encoding




Parametric Encoding (Dense Grid):

- Computationally efficient: few feature vectors are updated per sample
- Memory waste: too many features in **empty** space



# Multiresolution Hash Encoding

(a) No encoding



411 k + 0 parameters  
11:28 (mm:ss) / PSNR 18.56

(b) Frequency  
[Mildenhall et al. 2020]



438 k + 0  
12:45 / PSNR 22.90

(c) Dense grid  
Single resolution



10 k + 33.6 M  
1:09 / PSNR 22.35

(d) Dense grid  
Multi resolution



10 k + 16.3 M  
1:26 / PSNR 23.62

(e) Hash table (ours)  
 $T = 2^{14}$



10 k + 494 k  
1:48 / PSNR 22.61

(f) Hash table (ours)  
 $T = 2^{19}$

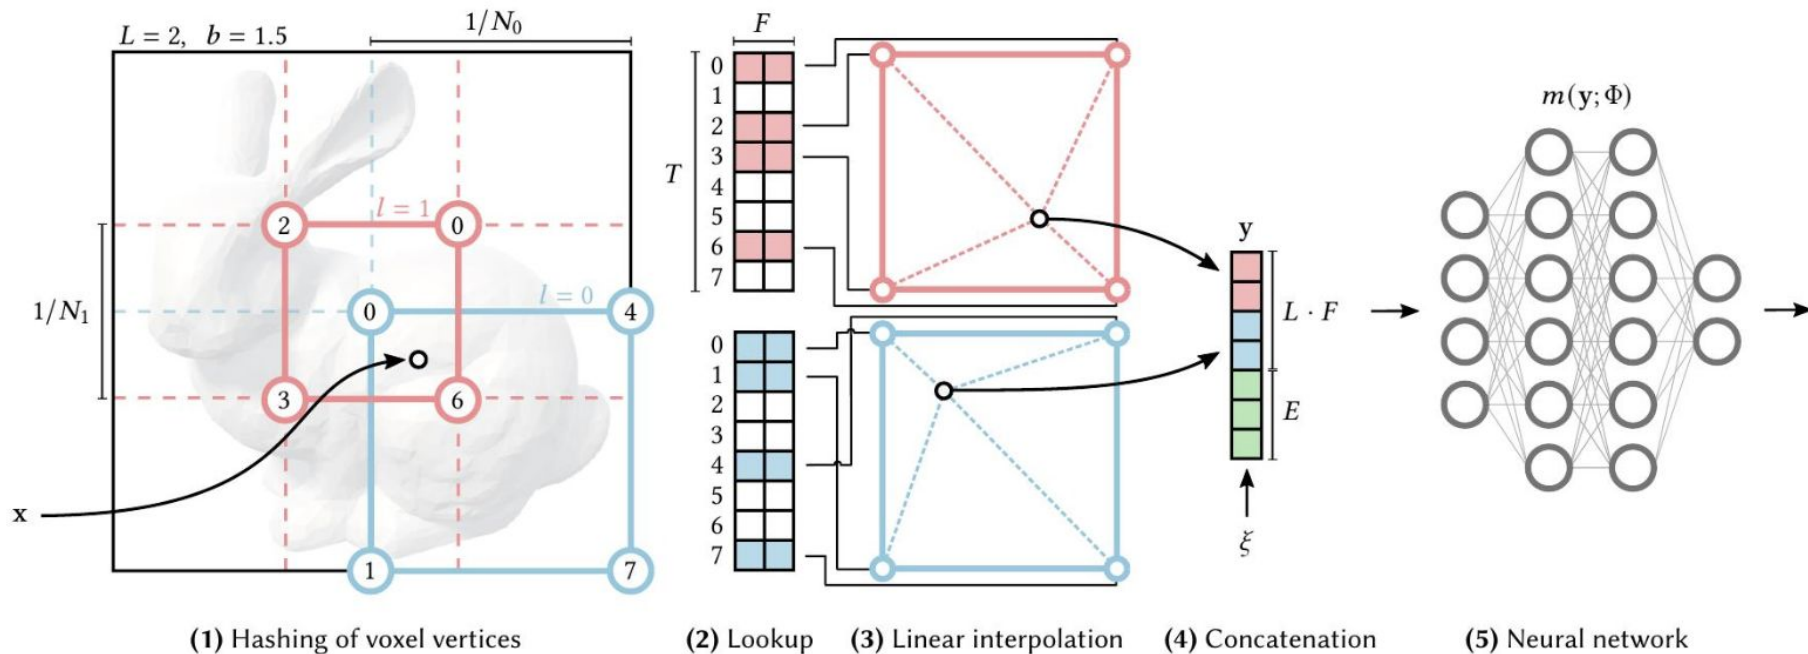


10 k + 12.6 M  
1:47 / PSNR 24.58

Parametric Encoding (Hash Table)

- Decrease

# Method - Overview



# Hash Encoding: Multi-Resolution

Neural network  $m(\mathbf{y}; \Phi)$  encodes inputs as  $\mathbf{y} = \text{enc}(\mathbf{x}; \theta)$  with trainable weights  $\Phi$  and encoding parameters  $\theta$ .

$\theta$  is split into  $L$  levels with up to  $T$  feature vectors of dimension  $F$ .

Grid resolutions range between  $[N_{\min}, N_{\max}]$  and are computed as:

$$N_l := \lfloor N_{\min} \cdot b^l \rfloor,$$
$$b := \exp\left(\frac{\ln N_{\max} - \ln N_{\min}}{L - 1}\right).$$



# Hash Encoding: Hash Function

For each level  $l$ , input  $\mathbf{x} \in \mathbb{R}^d$  is scaled to get  $\lfloor \mathbf{x}_l \rfloor$  and  $\lceil \mathbf{x}_l \rceil$ , which is associated with a voxel in  $\mathbb{Z}^d$ .

For coarse levels, mapping is **unique**. At finer levels, a hash function  $h : \mathbb{Z}^d \rightarrow \mathbb{Z}_T$  is employed.

The hash function is defined as:

$$h(\mathbf{x}) = \left( \bigoplus_{i=1}^d x_i \pi_i \right) \bmod T$$

$\pi_i$  are some large prime numbers. For each of the  $L$  levels, interpolated vectors are merged with auxiliary inputs (view direction or textures)  $\xi \in \mathbb{R}^E$  to form  $\mathbf{y} \in \mathbb{R}^{LF+E}$ .

# Hash Encoding: Collision Resistance



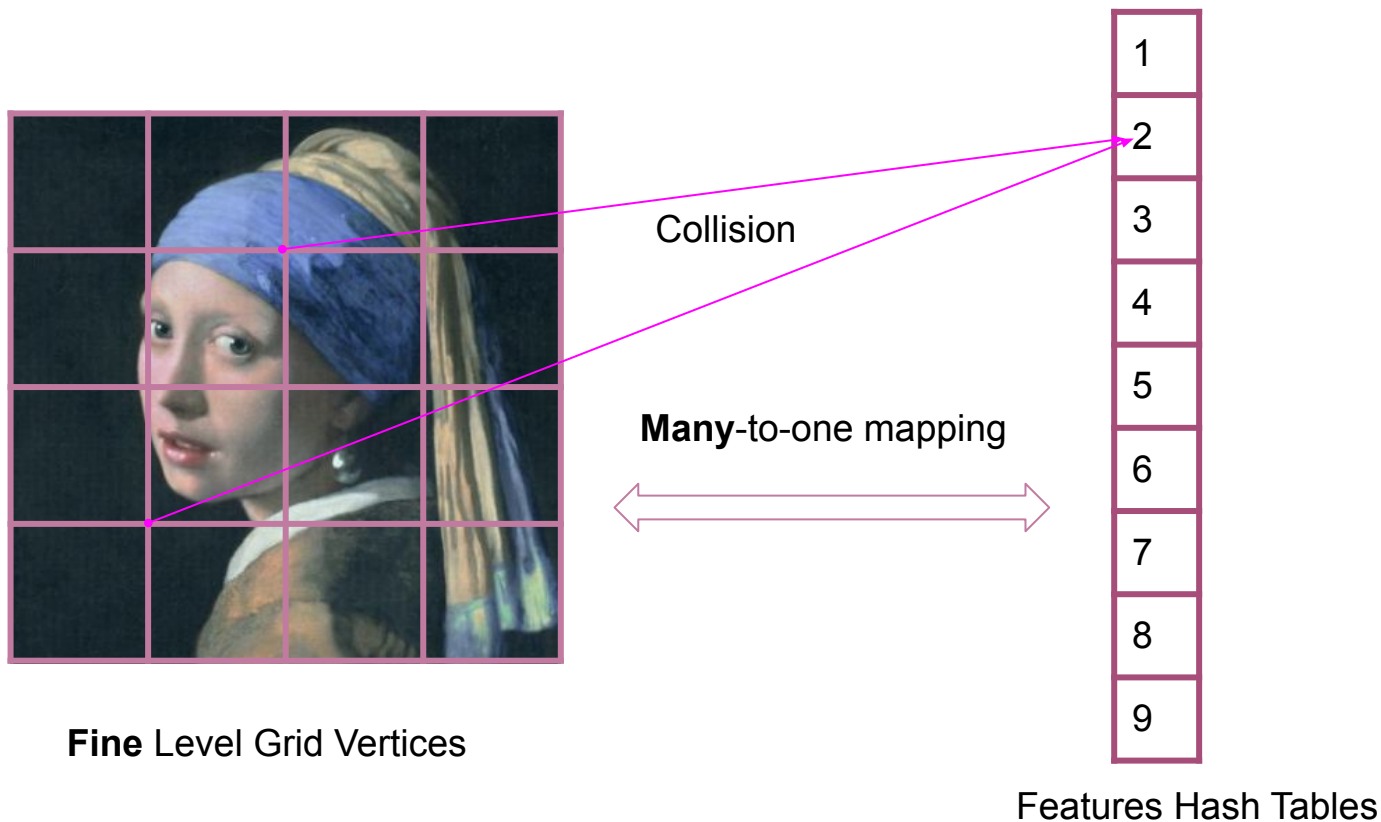
**Coarse** Level Grid Vertices

One-to-one mapping  
↔

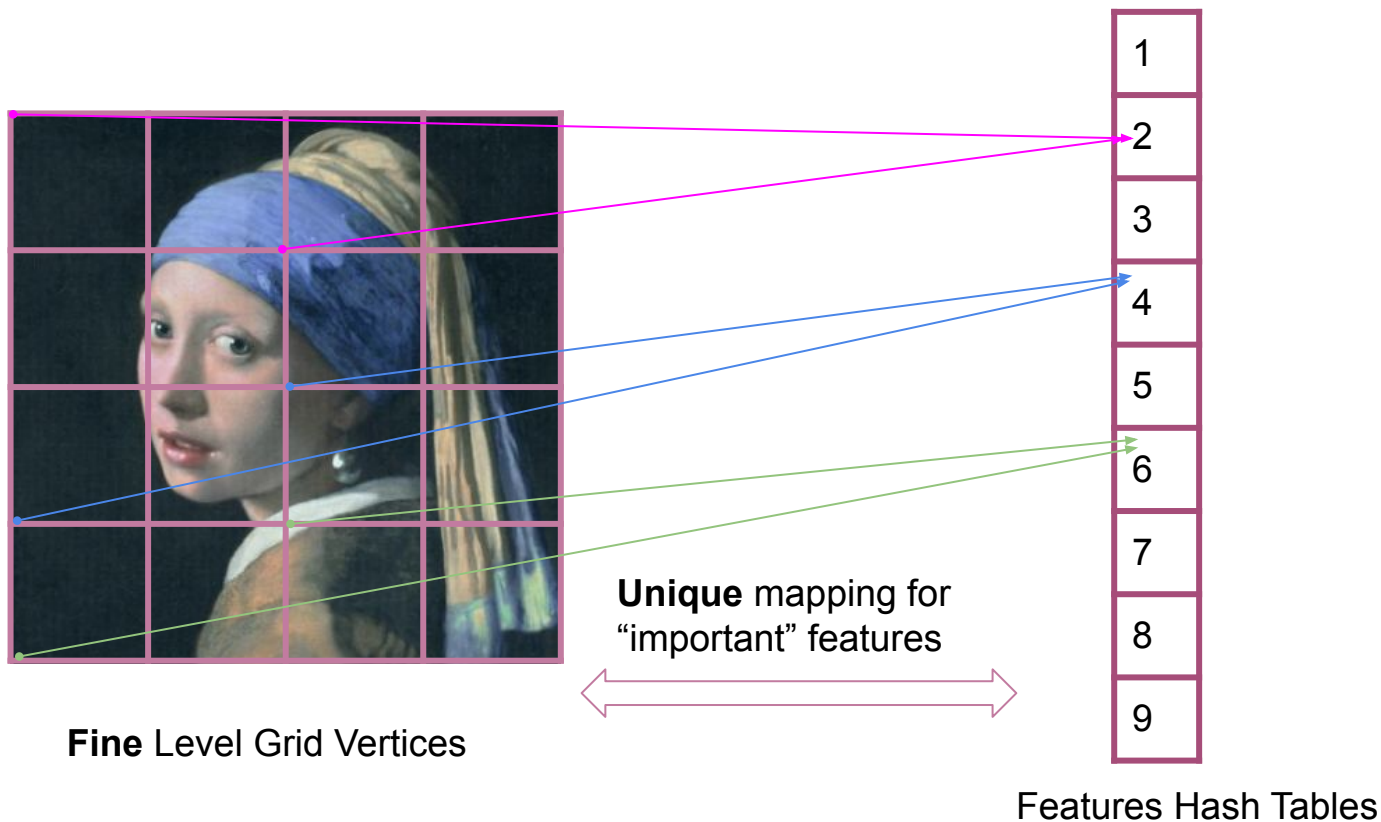
1
2
3
4
5
6
7
8
9

Features Hash Tables

# Hash Encoding: Collision Resistance



# Hash Encoding: Learnable Indices?

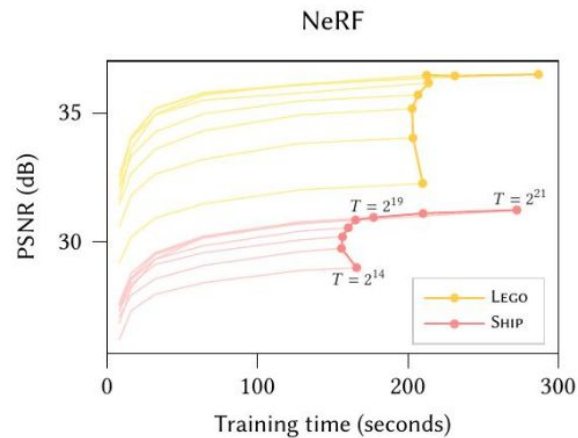
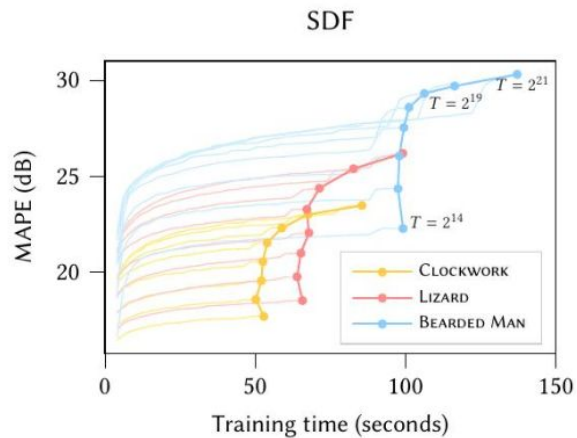
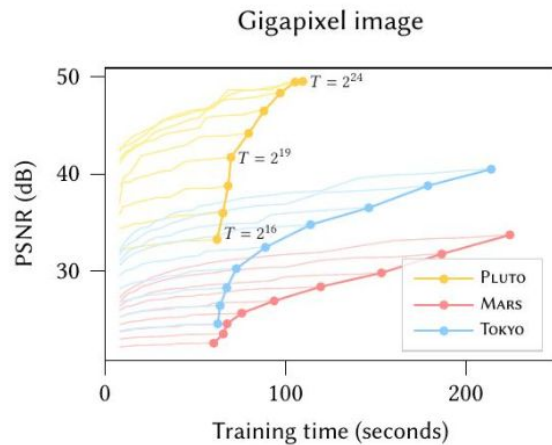


# Experiments

- (1) **Gigapixel image:** the MLP learns the mapping from 2D coordinates to RGB colors of a high-resolution image.
- (2) **Neural signed distance functions (SDF):** the MLP learns the mapping from 3D coordinates to the distance to a surface.
- (3) **Neural radiance caching (NRC):** the MLP learns the 5D light field of a given scene from a Monte Carlo path tracer.
- (4) **Neural radiance and density fields (NeRF):** the MLP learns the 3D density and 5D light field of a given scene from image observations and corresponding perspective transforms.

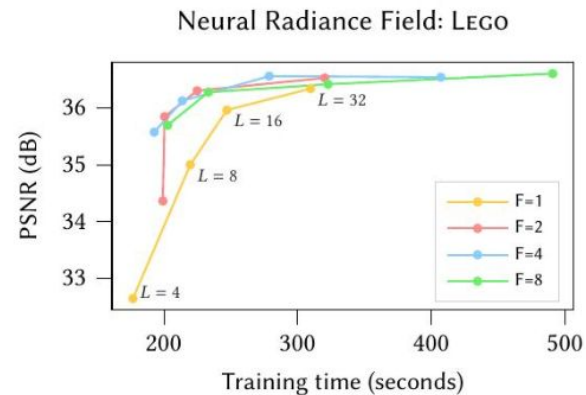
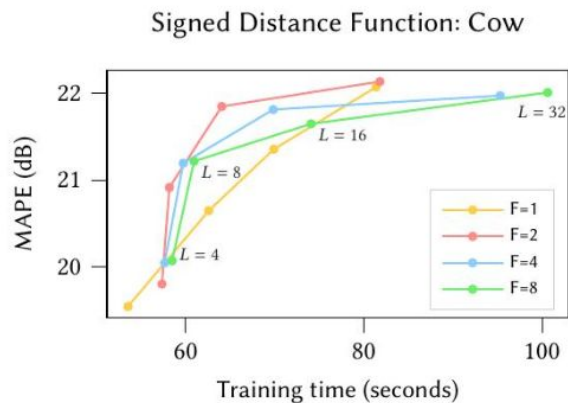
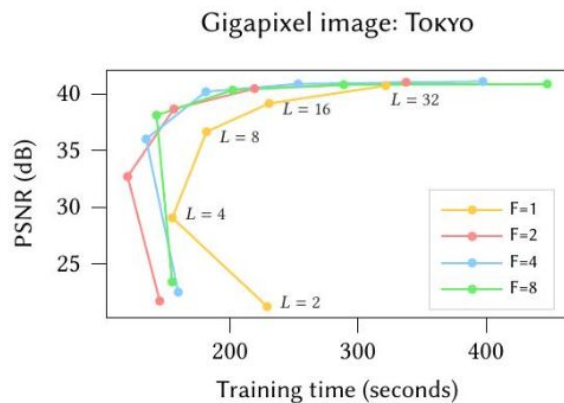


# Hyperparameters



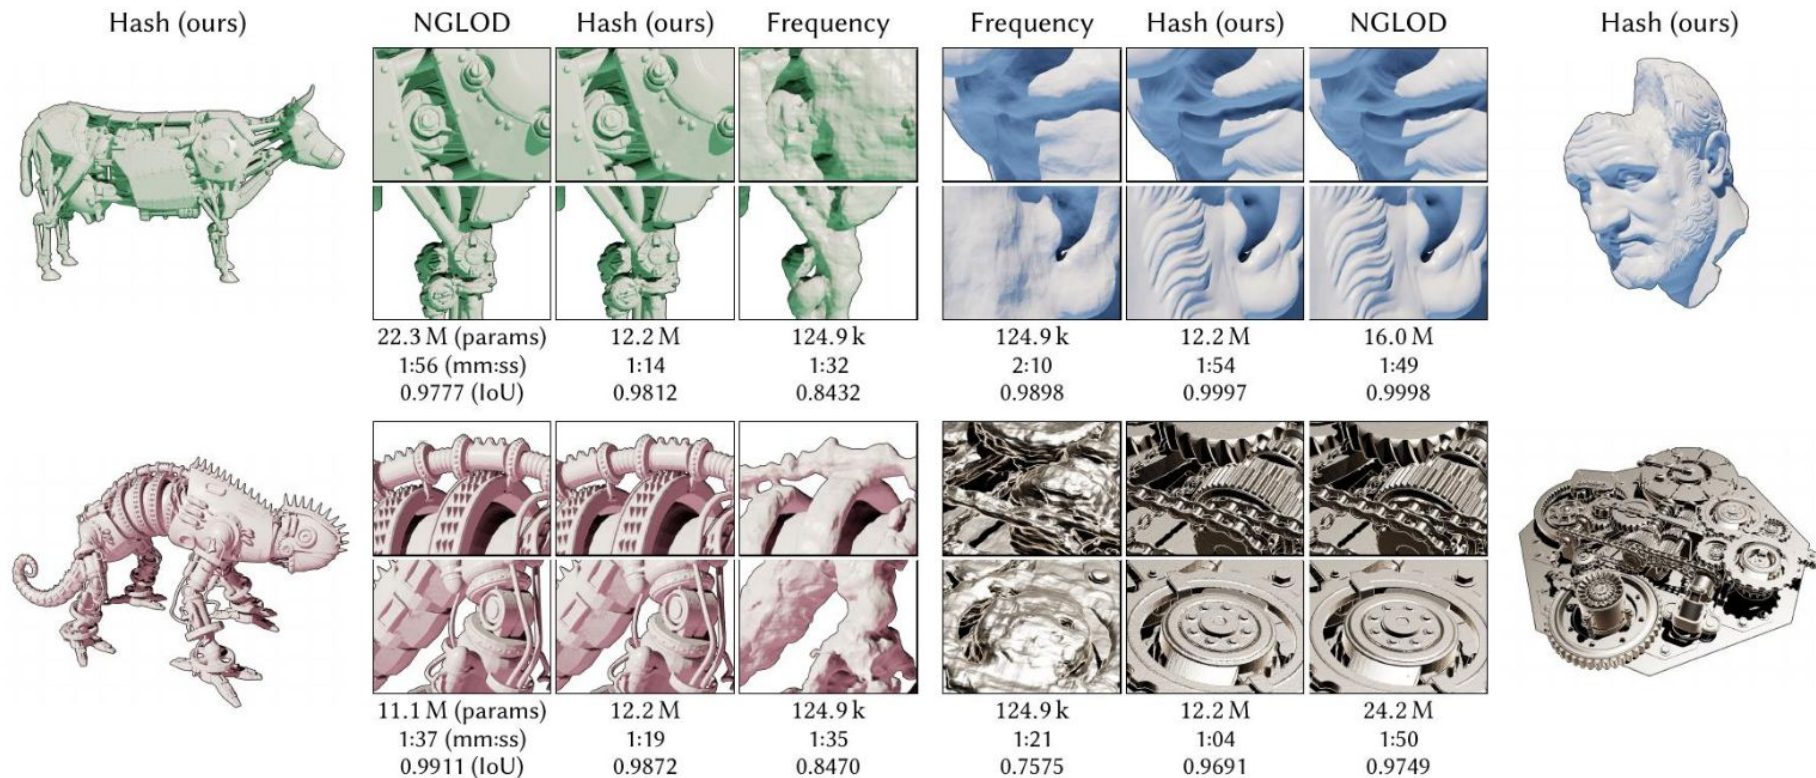
Varies Hash Table size  $T$

# Hyperparameters



Varies Feature Dimensionality **F**

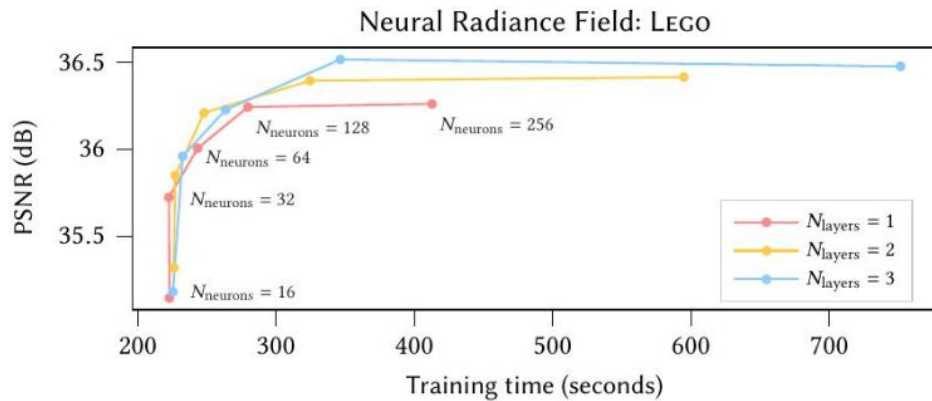
# Neural SDF




# NeRF

	Mic	FICUS	CHAIR	HOTDOG	MATERIALS	DRUMS	SHIP	LEGO	avg.
Ours: Hash (1 s)	26.09	21.30	21.55	21.63	22.07	17.76	20.38	18.83	21.202
Ours: Hash (5 s)	32.60	30.35	30.77	33.42	26.60	23.84	26.38	30.13	29.261
Ours: Hash (15 s)	34.76	32.26	32.95	35.56	28.25	25.23	28.56	33.68	31.407
Ours: Hash (1 min)	35.92 ●	33.05 ●	34.34 ●	36.78	29.33	25.82 ●	30.20 ●	35.63 ●	32.635 ●
Ours: Hash (5 min)	36.22 ●	33.51 ●	35.00 ●	37.40 ●	29.78 ●	26.02 ●	31.10 ●	36.39 ●	33.176 ●
mip-NeRF (~hours)	36.51 ●	33.29 ●	35.14 ●	37.48 ●	30.71 ●	25.48 ●	30.41 ●	35.70 ●	33.090 ●
NSVF (~hours)	34.27	31.23	33.19	37.14 ●	32.68 ●	25.18	27.93	32.29	31.739
NeRF (~hours)	32.91	30.13	33.00	36.18	29.62	25.01	28.65	32.54	31.005
Ours: Frequency (5 min)	31.89	28.74	31.02	34.86	28.93	24.18	28.06	32.77	30.056
Ours: Frequency (1 min)	26.62	24.72	28.51	32.61	26.36	21.33	24.32	28.88	26.669

## Size of the MLPs



## Hash Encoding enables a small MLP



## Linear model v.s. MLP



# References

Instant-NGP: <https://nvlabs.github.io/instant-ngp/assets/mueller2022instant.pdf>

Instant-NGP snapshot: [https://docs.nerf.studio/en/latest/nerfology/methods/instant\\_ngp.html](https://docs.nerf.studio/en/latest/nerfology/methods/instant_ngp.html)

NSVF: <https://arxiv.org/abs/2007.11571>

Neural Radiance Caching:

[https://research.nvidia.com/publication/2021-06\\_real-time-neural-radiance-caching-path-tracing](https://research.nvidia.com/publication/2021-06_real-time-neural-radiance-caching-path-tracing)

Thank You!